

28-29-30//NOVEMBRE 2017

OverNet
EDUCATION

La più importante
conferenza italiana
sulle tecnologie
Microsoft



28-29-30//NOVEMBRE
WPC
M I L A N O **2017**

[Graph DB] in SQL Server 2017, Azure SQLDB e Azure Cosmos DB : da 0 a 100 in 60 minuti

IGOR PAGLIAI

CLOUD AND DATA ARCHITECT

CSE INDUSTRY TEAM

MICROSOFT CORP.

 @igorpag



Agenda

- Introduzione ai “Grafì” - 101
 - Quali tecnologie
 - Modello relazionale vs. grafi
 - Scenari e concetti di base
- SQL Server 2017 & Azure SQLDB
 - Supporto e funzionalità
 - Limitazioni e «roadmap»
- Azure Cosmos DB
 - Caratteristiche e funzionalità
 - Supporto specifico ai grafi
 - Gremlin!
- Comparazione
 - PRO e CONTRO di ogni soluzione
- Q&A

OverNet
EDUCATION





Introduzione ai Grafi



Eulero, [Il problema dei ponti di Königsberg](#)

Rise of graph databases

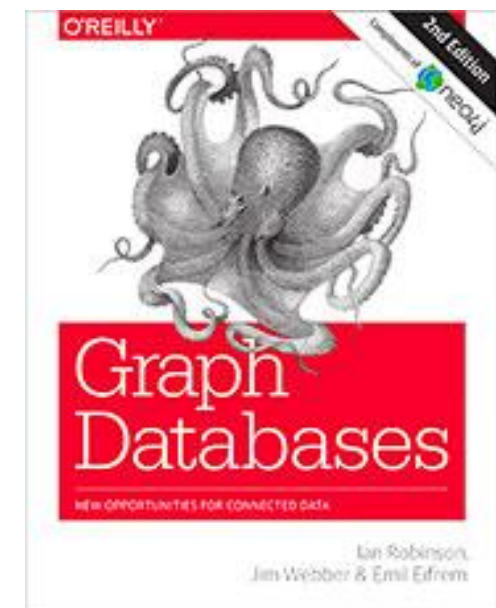
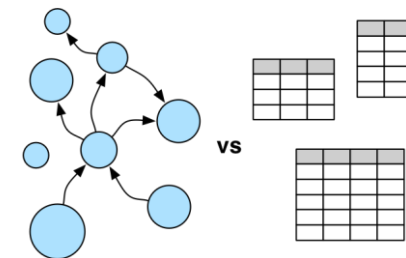
- Rising in popularity: why?
- Gartner: Leveraging 5 graphs provides *sustained competitive advantage*
 - social, intent, consumption, interest, and mobile
- Major database vendors
 - Neo4j Inc.
 - Azure Cosmos DB (Microsoft ☺)
 - OrientDB
 - Titan + Cassandra
 - Giraph (Facebook)
- Industry leaning towards multi-model

27 systems in ranking, November 2017

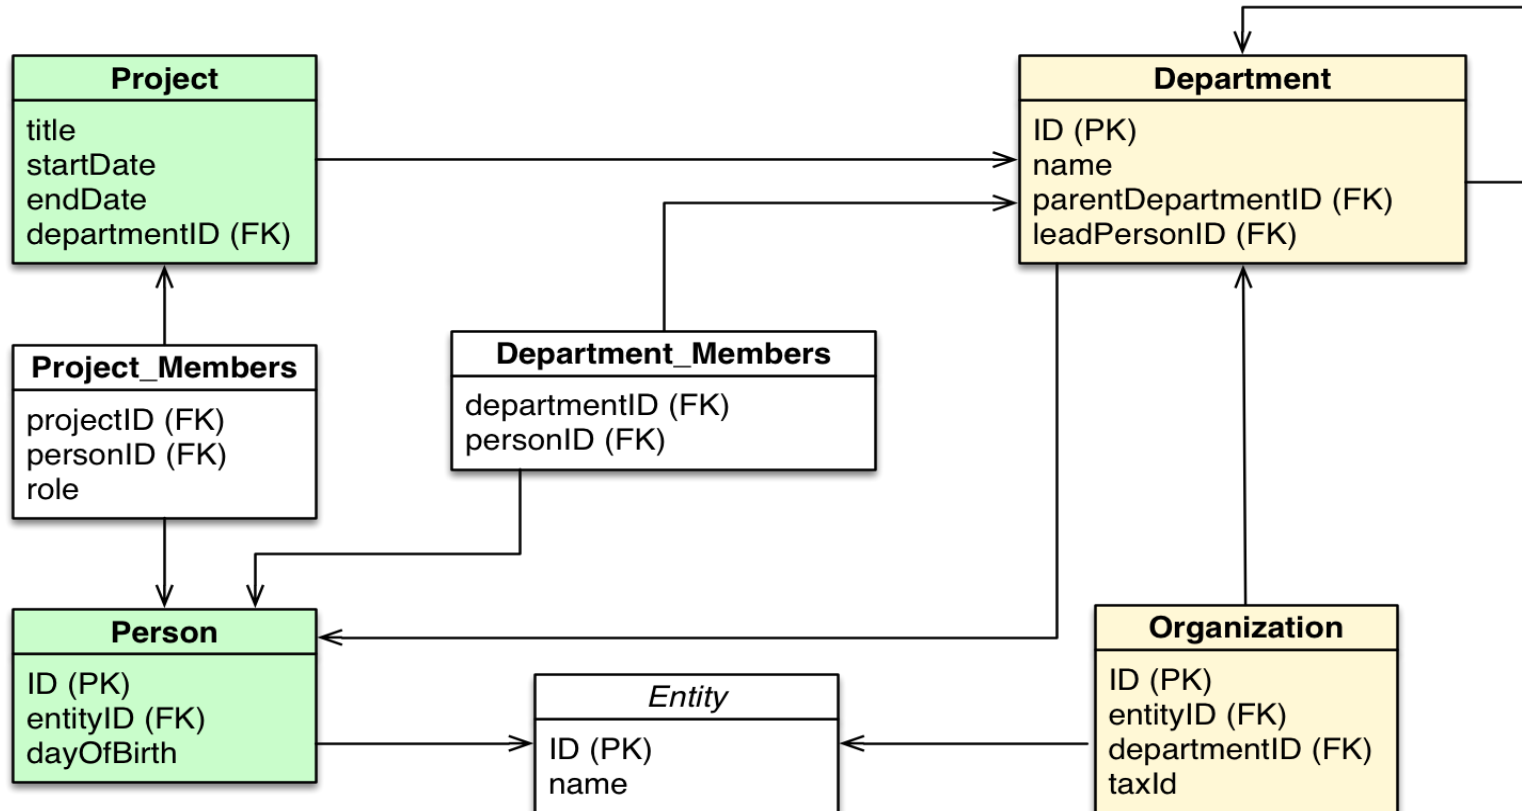
Rank			DBMS	Database Model	Score		
Nov 2017	Oct 2017	Nov 2016			Nov 2017	Oct 2017	Nov 2016
1.	1.	1.	Neo4j +	Graph DBMS	38.45	+0.50	+1.70
2.	2.	↑ 4.	Microsoft Azure Cosmos DB +	Multi-model i	13.03	+0.40	+9.78
3.	3.	↓ 2.	OrientDB +	Multi-model i	6.10	-0.03	+0.03
4.	4.	↓ 3.	Titan	Graph DBMS	5.68	+0.21	+0.22
5.	5.	↑ 6.	ArangoDB	Multi-model i	3.15	-0.01	+0.87
6.	6.	↓ 5.	Virtuoso	Multi-model i	1.88	+0.03	-0.67
7.	7.	7.	Giraph	Graph DBMS	1.02	-0.01	+0.05
8.	8.	↑ 12.	GraphDB +	Multi-model i	0.74	+0.10	+0.56
9.	9.	9.	AllegroGraph +	Multi-model i	0.60	-0.02	+0.12
10.	10.	↓ 8.	Stardog	Multi-model i	0.54	+0.01	+0.01


Neo4j

- Most famous and widely adopted
 - Graph model only, no multi-model, OSS, Java
- Cypher primary language, Gremlin also supported
- Interesting quote from author: “*Relational models lacks relations!*”
 - Relational model is not natural and force logical remapping
 - M:M relationships, Normal Forms, etc.
- Native “Compute” and “Storage” layers, no hybrid
 - “Index-free Adjacency”
- Native algorithm support
 - Transitive closure, arbitrary length traversal, etc.
- Super efficient storage layer
 - Fixed length records: offset calculation by ObjectID
 - Separate files for Edge and Nodes
 - External storage for blobs/binary
- No PaaS or SaaS (so far)
- No global distribution (so far)
- No horizontal partitioning/sharding (so far)



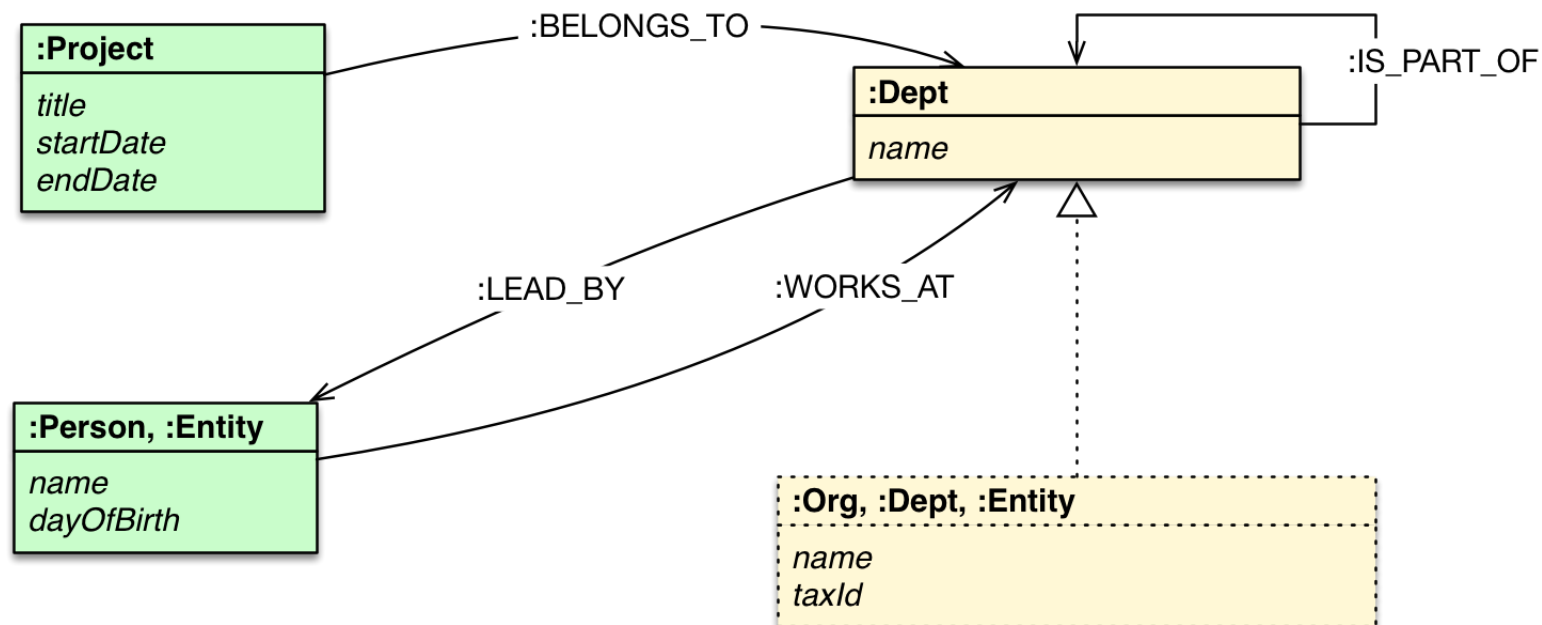
Relational vs. Graph: Models



- 1NF – First normal form
- 2NF – Second normal form
- 3NF – Third normal form
- EKNF – Elementary key normal form
- BCNF – Boyce–Codd normal form
- 4NF – Fourth normal form
- ETNF – Essential tuple normal form 
- 5NF – Fifth normal form
- 6NF – Sixth normal form
- DKNF – Domain/key normal form

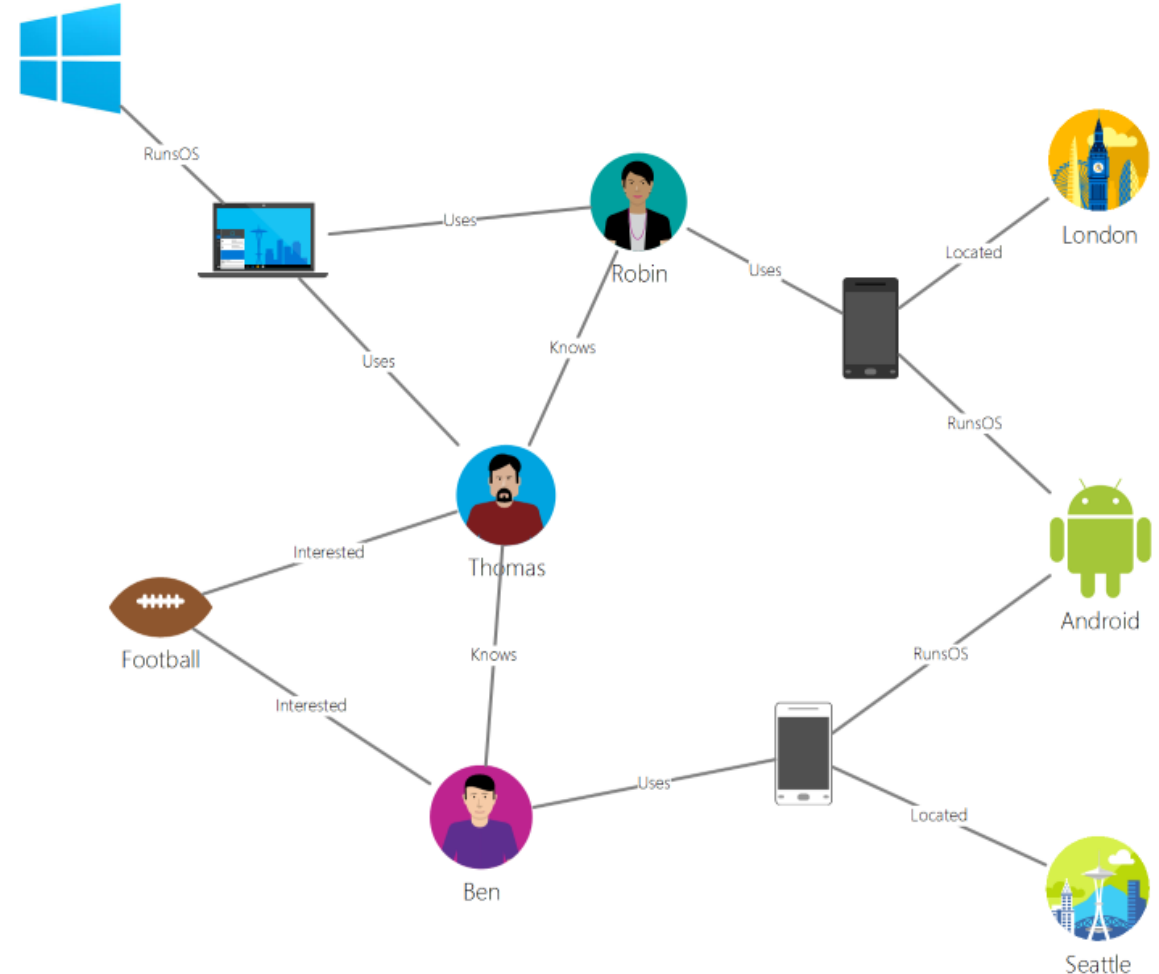
Source: <https://neo4j.com/developer/graph-db-vs-rdbms>

Relational vs. Graph: Models



Why use graph databases?

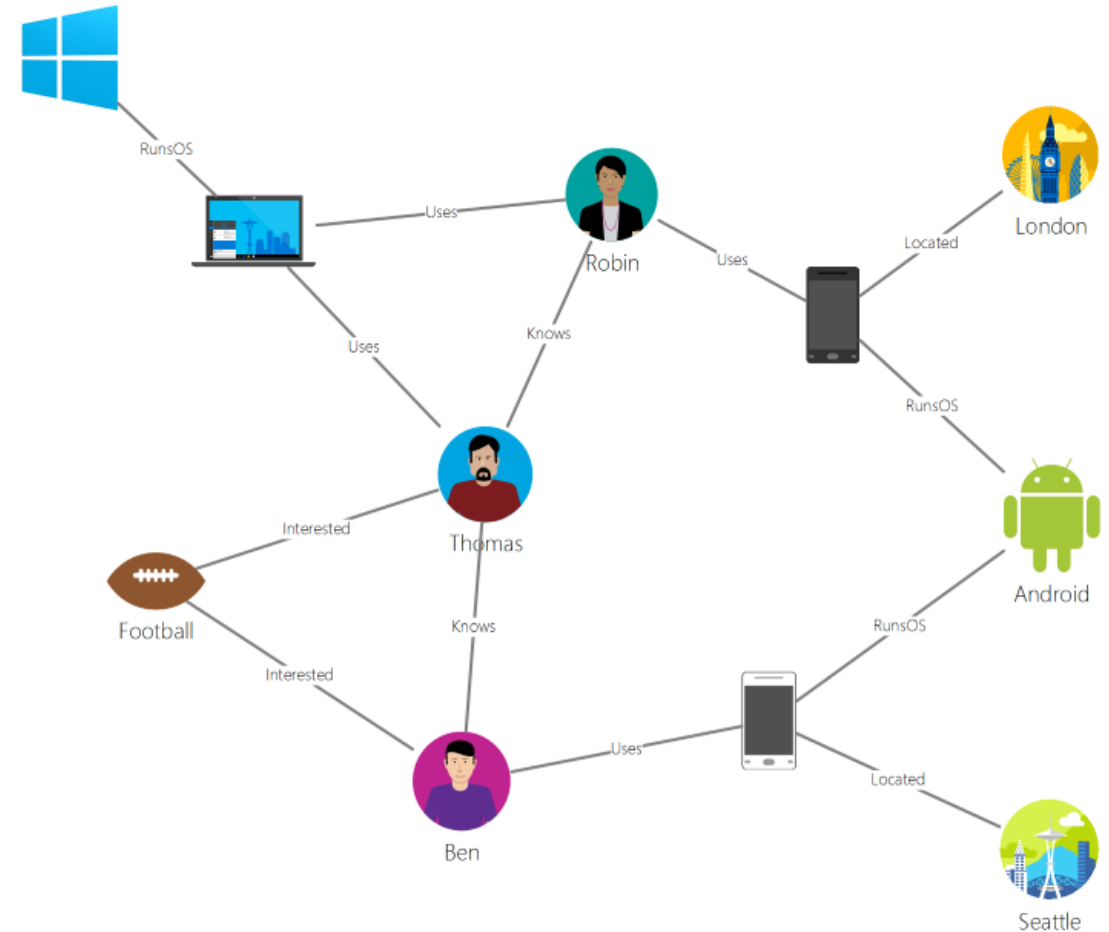
- Natural and intuitive way to model the real world
- Graph traversal/navigation
- Relationships are first-class concepts
- No Joins!
- Built for fast graph computations
- Graph databases are traditionally flexible, No Schema (NoSQL family)



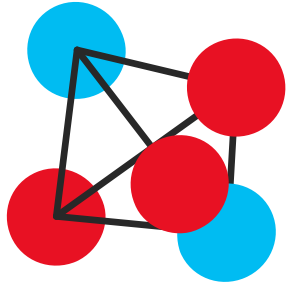
Graph database concepts

Vertices, Edges and Properties

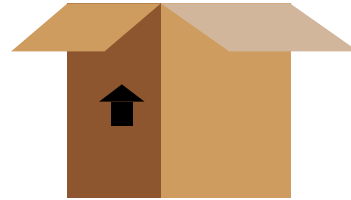
- Vertex (or Node) = any entity
- Edge = relationship, connects two vertices
- Vertex and edge have “Labels” and “Properties” (Key-Value)
- Graph = $G(V,E)$ = Collection of Vertices and Edges
- Traversals = query over graphs



Graph database scenarios



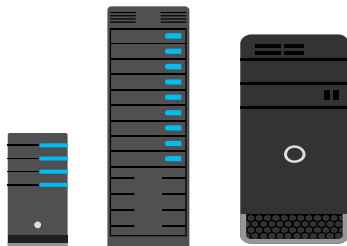
Social Networking



Recommendations



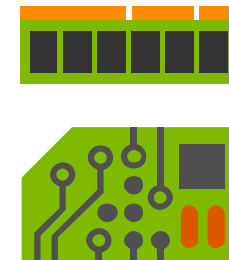
Security &
Fraud Detection



Network & Operations



Customer 360



IoT (Logistics)

SQL Server 2017 & Azure SQLDB

GRAPH PROCESSING WITH SQL SERVER 2017 & AZURE
SQL DATABASE

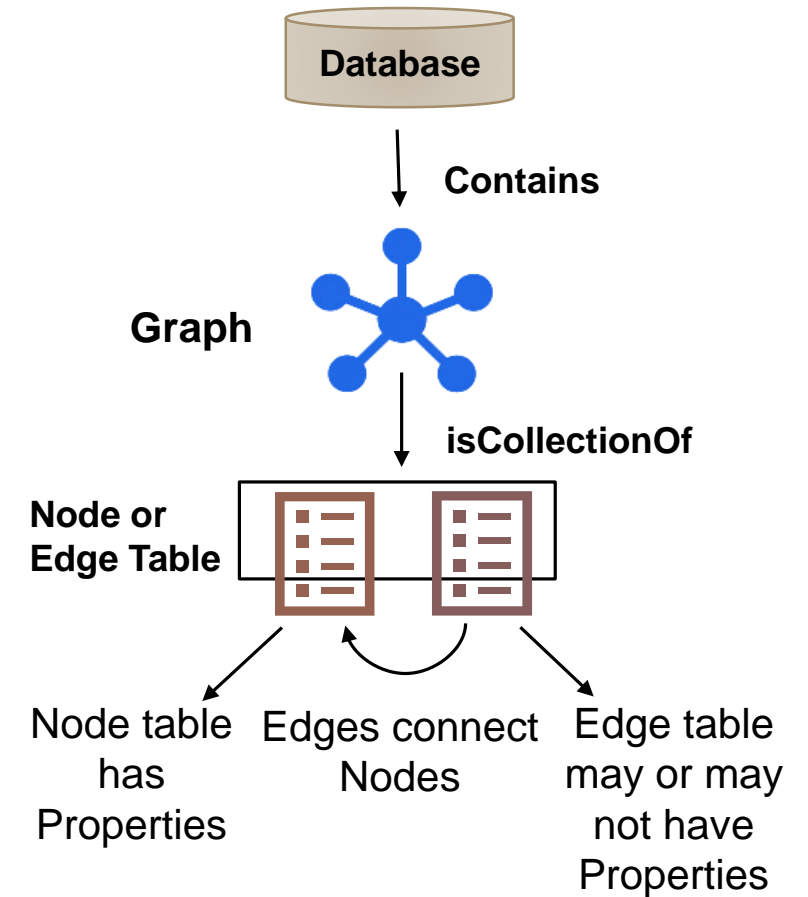
Introducing SQL Graph

A collection of Node and Edge tables in the database

Language Extensions

- DDL Extensions – create node/edge tables
- Query Language Extensions – New built-in: **MATCH**, to support pattern matching and traversals

Tooling and Eco-system



DDL Extensions

Create Nodes and Edges

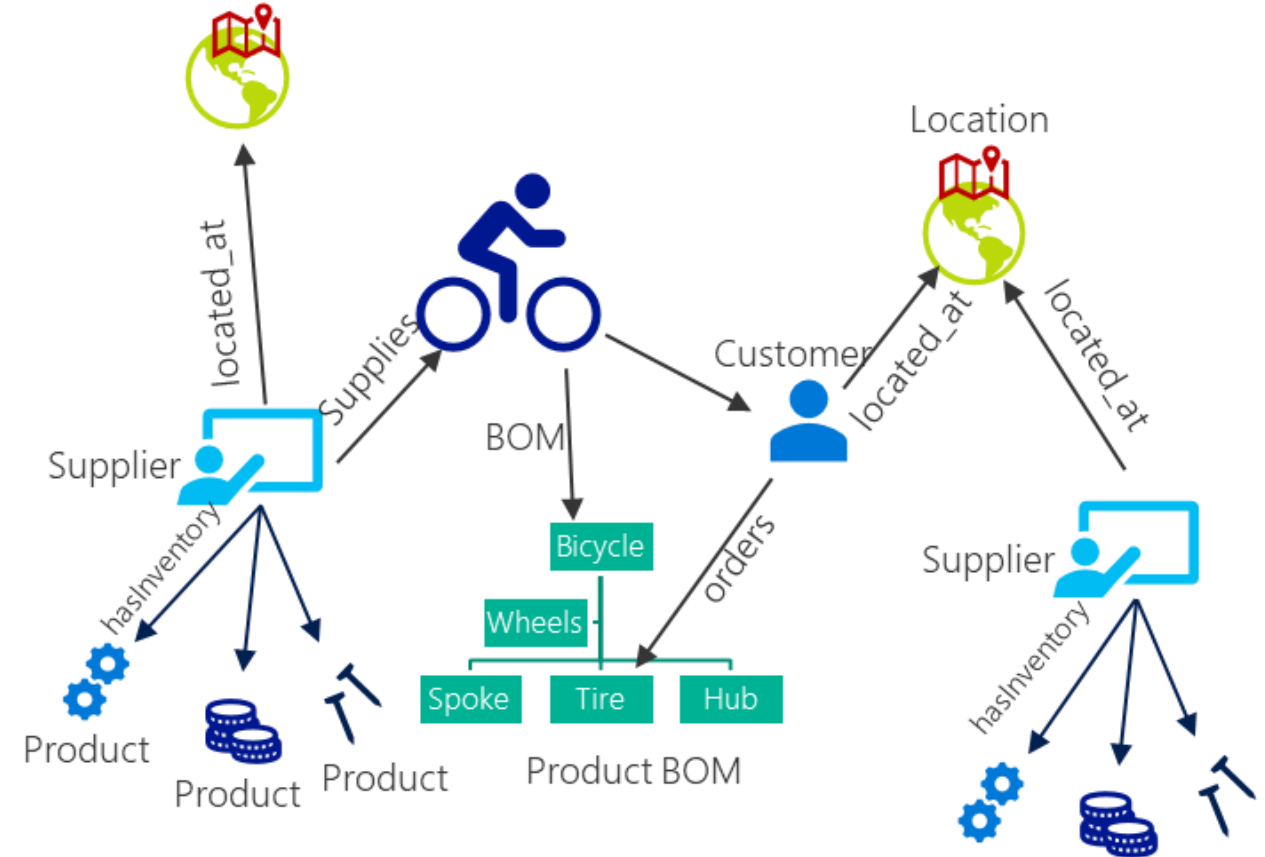
Properties associated with Nodes and Edges

```
CREATE TABLE Product (ID INTEGER
PRIMARY KEY, name VARCHAR(100)) AS
NODE;
```

```
CREATE TABLE Supplier (ID INTEGER
PRIMARY KEY, name VARCHAR(100)) AS
NODE;
```

```
CREATE TABLE hasInventory AS EDGE;
```

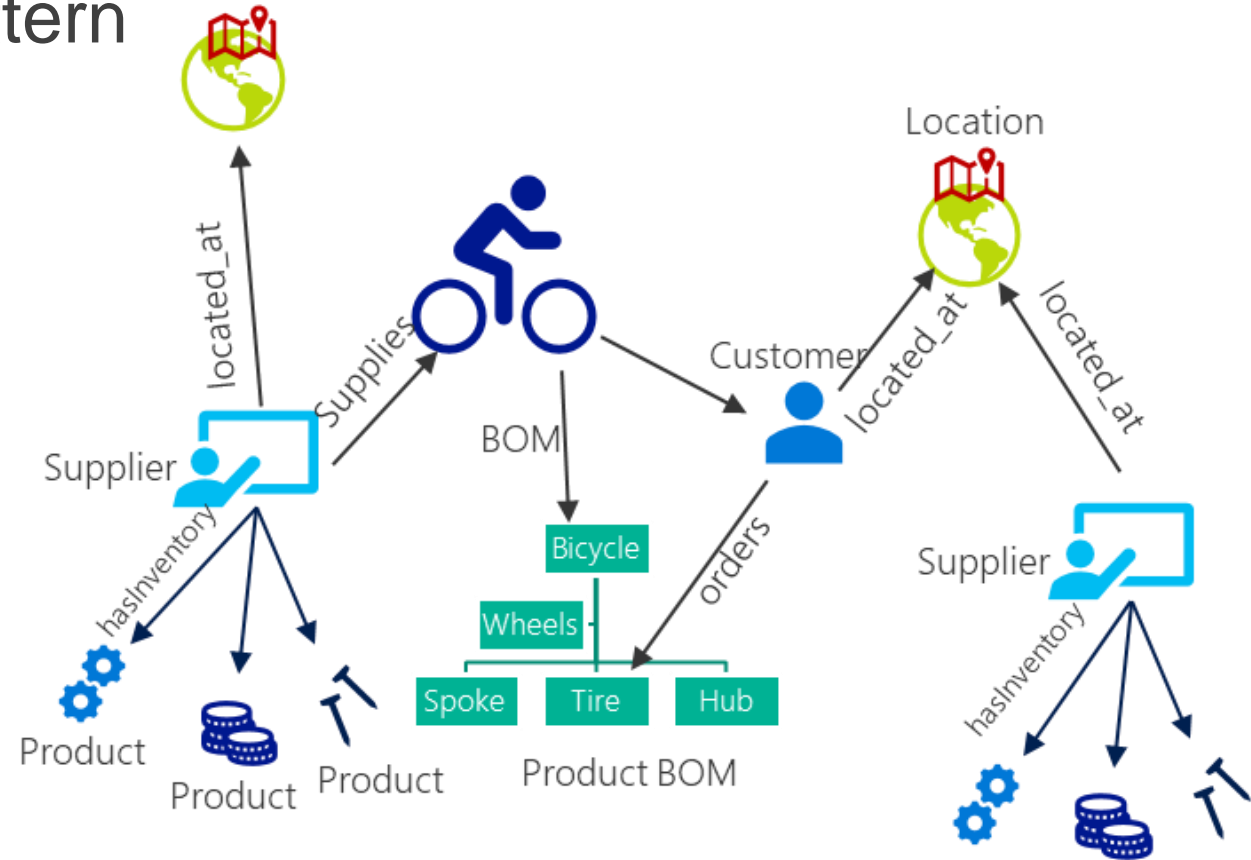
```
CREATE TABLE located_at(address
varchar(100)) AS EDGE;
```



Query Language Extensions

Multi-hop navigation and join-free pattern matching using **MATCH** predicate:

```
SELECT Prod.name as ProductName,
       Sup.name as SupplierName
FROM Product Prod, Supplier Sup,
     hasInventory hasIn,
     located_at supp_loc,
     Customer Cus,
     located_at cust_loc,
     orders, location loc
WHERE
MATCH(
  cus-(orders)->Prod<-(hasIn)-Sup
)
```



Integrated into the SQL Engine

Existing tools all work out of the box, for example SSMS, backup and restore, import / export, etc

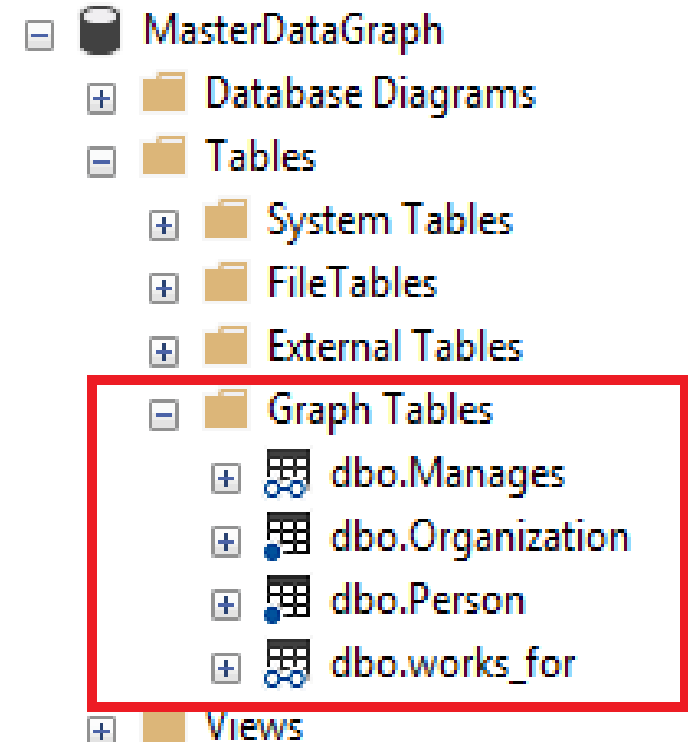
Queries can join existing tables and graph node / edge tables

Columnstore Indexes, ML (R / Python), HA etc.

Security and Compliance: Dynamic Data Masking, Row Level Security

Available on SQL Server on Linux as well!

NEW tool: SQL Operations Studio (SQLOpsStudio) for Windows, Linux and MacOS



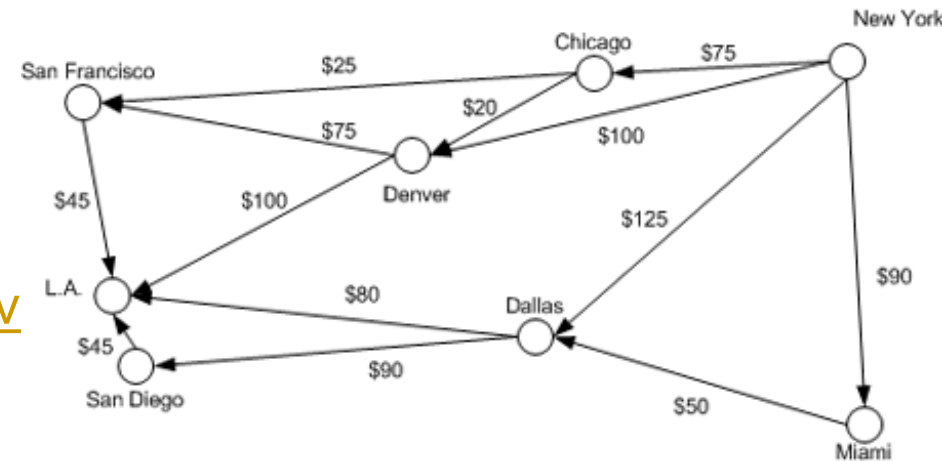
Relational vs. Graph

- Graph and relational designs can answer the same questions
 - But if traversal of relationships define the primary application requirements, Graph can solve this more intuitively and with less code

- Shortest Paths

- Find all paths from a given node to another node
- Find shortest paths to all other (reachable) nodes
- Sample T-SQL implementation in https://github.com/arvindshmicrosoft/YelpDatasetSQLServer/3_GraphAlgorithms.sql
- Gremlin Example:

```
g.V('robin')  
  .repeat(out())  
  .until(has('id', 'ben'))  
  .path()
```



Why SQL Graph over 100s of other Graph DBs?

SQL Server and SQL DB as the main Data Hub

1. Bring computation closer to data: Machine Learning & AI, Blobs Hadoop, Graph
2. Handle all data formats : relational, FileStream, XML, JSON
3. Deployment flexibility: Linux, Windows, Docker & Containers, Azure, On-Premises
4. Enterprise grade features: security, audit, encryption, HA&DR, performances, etc.
5. Familiar: continue using existing skills – no surprises!

Current limitations

- Future Improvements
 - Shortest Path native implementation
 - Arbitrary length traversals:
Find friends 1-5 hops away

MATCH (Person-(isFriendOf*1:5)->Person)

- Heterogeneous associations
 - Find Sessions or Speakers that one likes
- Language enhancements – MERGE DML
 - Create if not exist
- Edge Constraints
 - Edges can connect only to specific Nodes



DEMO

Using Graph DB support in SQL Server 2017 and Azure SQLDB

Azure Cosmos DB

Linked Session: DEV006 - CosmosDB - La nuova frontiera del BigData e NoSql
(Today in Sala Verde at 5:00pm)



Bring your data (and API) as it is



Apache Cassandra



Table API

MongoDB API



Key-value



Column-family



Document



Graph

Elastic scale out
of storage & throughput

Guaranteed low latency at the 99th percentile

Five well-defined consistency models

Turnkey global distribution

Comprehensive SLAs



Azure Cosmos DB Graph API



Latency & Consistency Models

Guaranteed low latency at P99 (99th percentile)

Requests are served from local region

Single-digit millisecond latency worldwide

Write optimized, latch-free database engine designed for SSD

Synchronous automatic indexing at sustained ingestion rates



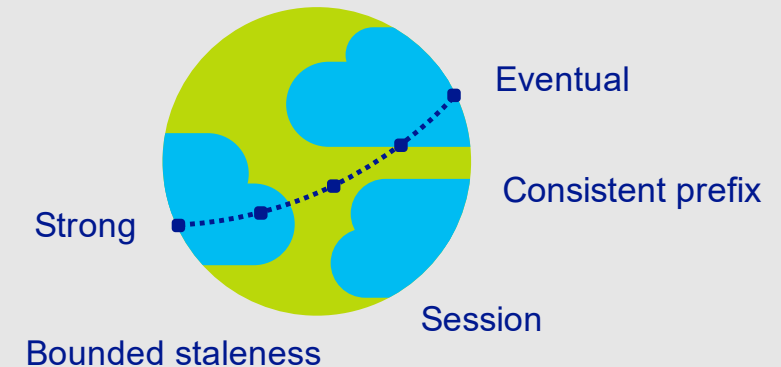
Multiple, well-defined consistency choices

Global distribution forces us to navigate the CAP theorem

Writing correct distributed applications is hard

Five well-defined consistency levels

Programmatically change at anytime, can be overridden on a per-request basis



Scalability & Indexing

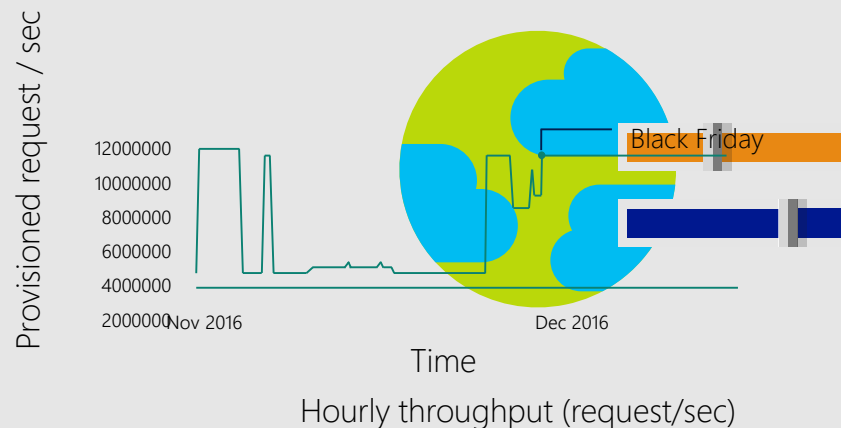
Elastically scalable storage and throughput (RU)

Single machine is never a bottle neck, transparent server-side partition management

Elastically scale storage (GB to PB) and throughput (100 to 100M req/sec) across many machines and multiple regions

Automatic expiration via policy based TTL

Pay by the hour, change throughput at any time for only what you need



Schema-agnostic, automatic indexing

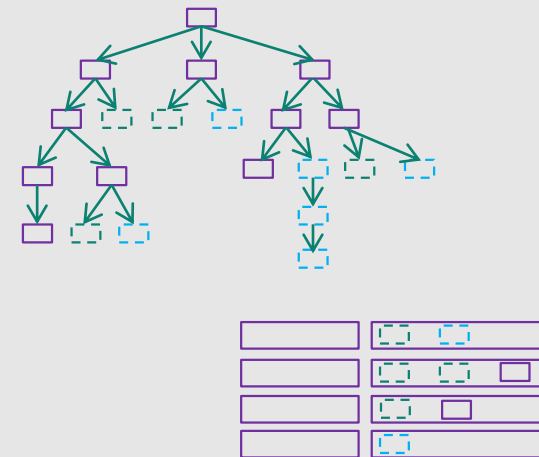
At global scale, schema/index management is painful

Automatic and synchronous indexing

Hash, range, and geospatial

Works across every data model

Highly write-optimized database engine



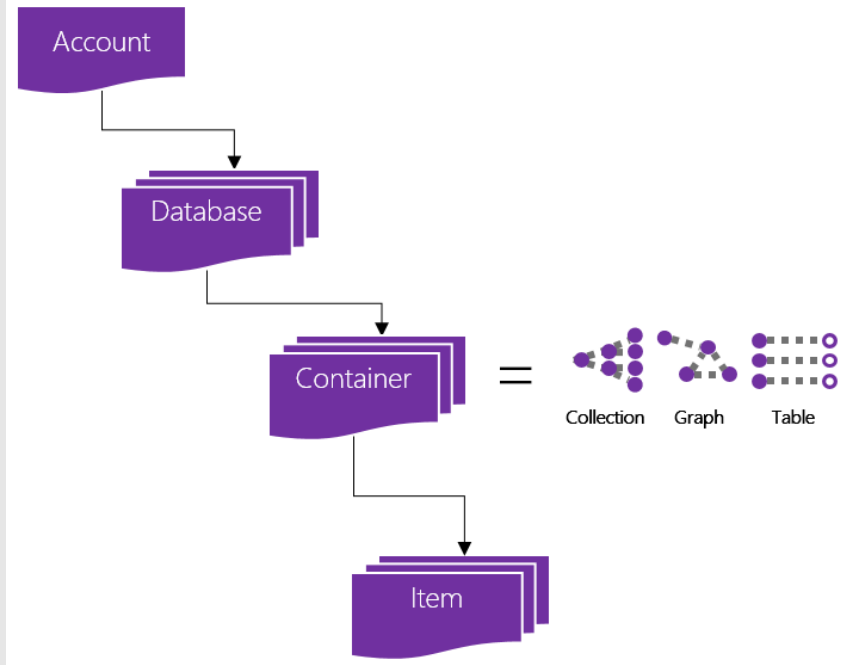
Partitioning

Resource Model

Resources identified by their logical and stable URI

Hierarchical overlay over horizontally partitioned entities; spanning machines, clusters and regions

Extensible custom projections based on specific type of API interface



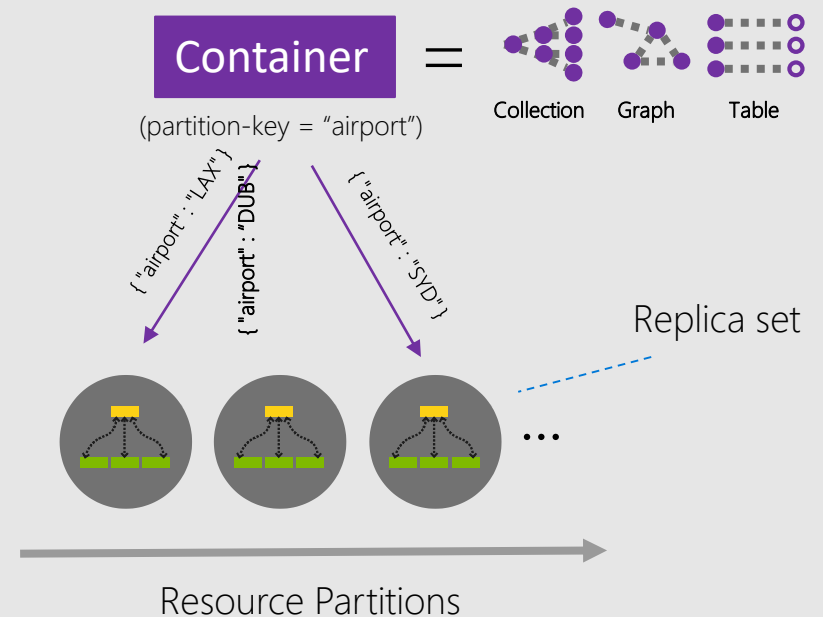
Horizontal Partitioning

Containers are horizontally partitioned

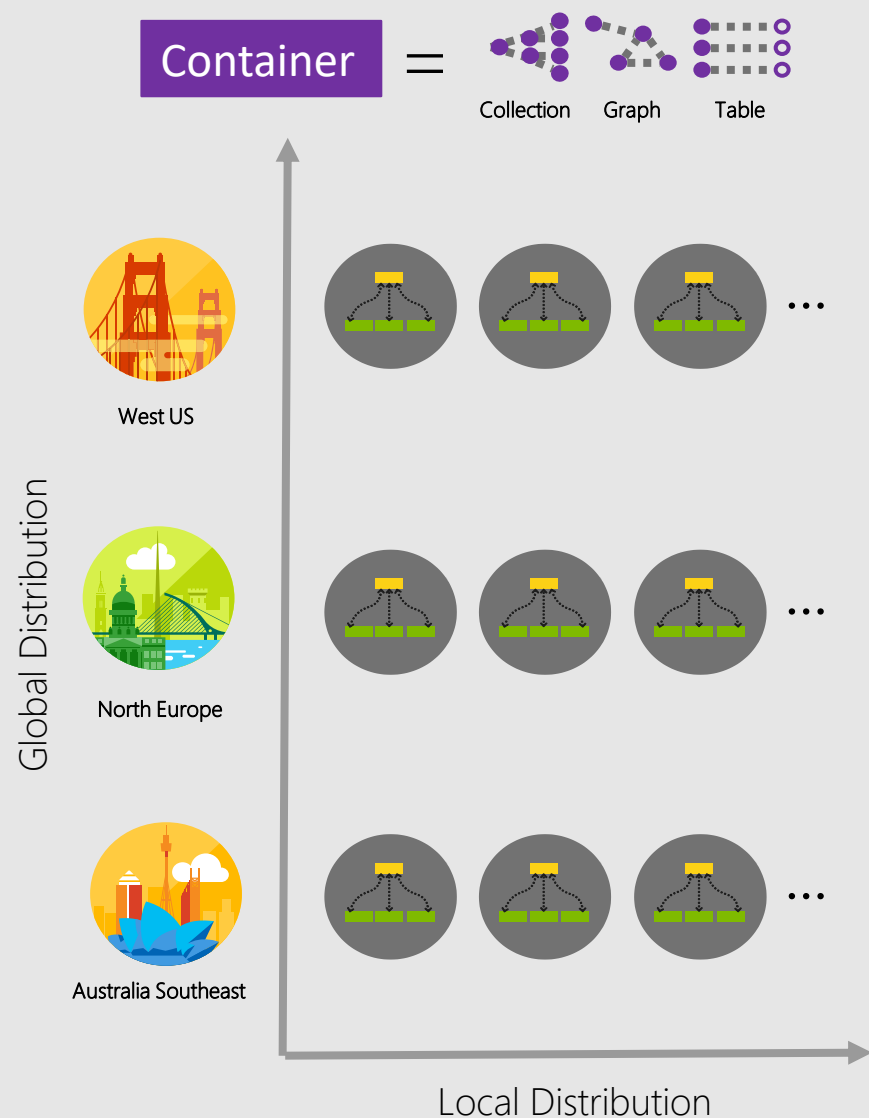
Each partition made highly available via a replica set

Partition management is transparent and highly responsive

Partitioning scheme is dictated by a "partition-key"



Global Distribution



Global Distribution

All resources are horizontally partitioned and vertically distributed

Distribution can be within a cluster, x-cluster, x-DC or x-region

Replication topology is dynamic based on consistency level and network conditions

The service handles routing query requests to the right partition using the partition key

Best Practices: Partitioning

All items with the same partition key will be stored in the same partition

Multiple partition keys may share the same partition using hash-based partitioning

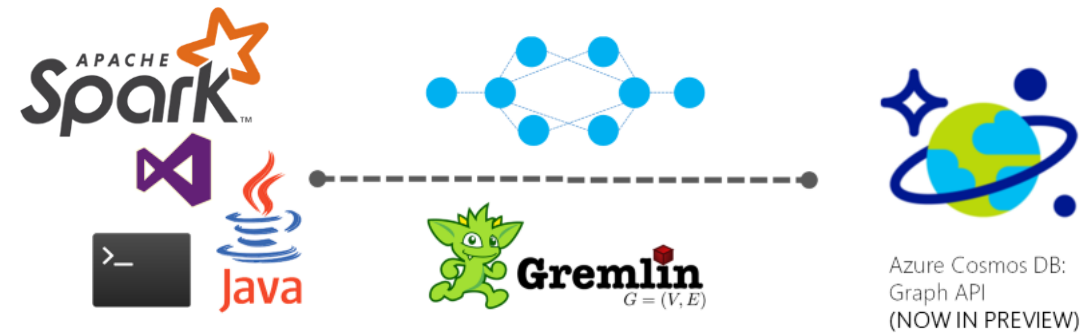
Select a partition key which provides even distribution of storage and throughput (req/sec) at any given time to avoid storage and performance bottlenecks

Partition key should be represented in the bulk of queries for read heavy scenarios to avoid excessive fan-out.

Partition key is the boundary for cross item transactions. Select a partition key which can be a transaction scope.

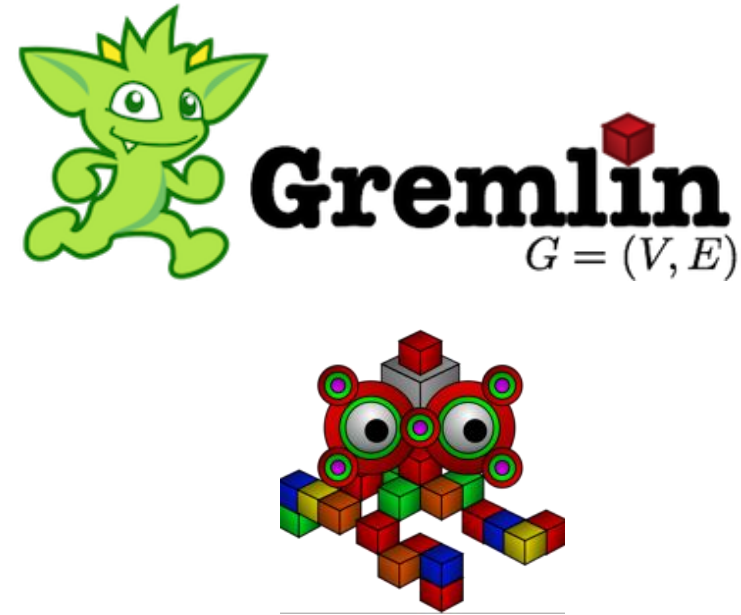
APIs and programmability: Gremlin!

- Part of **TinkerPop** Apache Project
 - Graph traversal language & API
 - Console **Client** and **Server**
- Work with any **TinkerPop** client driver
 - Java, Python, Node.js, etc.
 - GraphSON and Gryo wire formats
- First-party .NET client SDK
- Direct connectivity to CosmosDB
 - Also supported in Titan, DSE Graph, Neo4j
- Gremlin ecosystem just works
 - Work with Apache Spark on HDInsight for graph analytics



Developing with Gremlin 101 (Demo)

- Using any Gremlin driver
- Add a vertex
- Add an edge
- Filter by a property value
- Traverse the graph (join)
- More complicated examples
- From SQL to Gremlin



```
C:\WINDOWS\system32\cmd.exe

  \,,,/
  (o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin>
```



DEMO

Using GREMLIN API to access Azure Cosmos DB for Graph traversal

The Azure Cosmos DB Graph API (Gremlin) service will move out of preview and into general availability **by the end of the year 2017**

feedback on what you want to see next

Developer Forums on Stack Overflow



Comparison Recap

- **Neo4j**

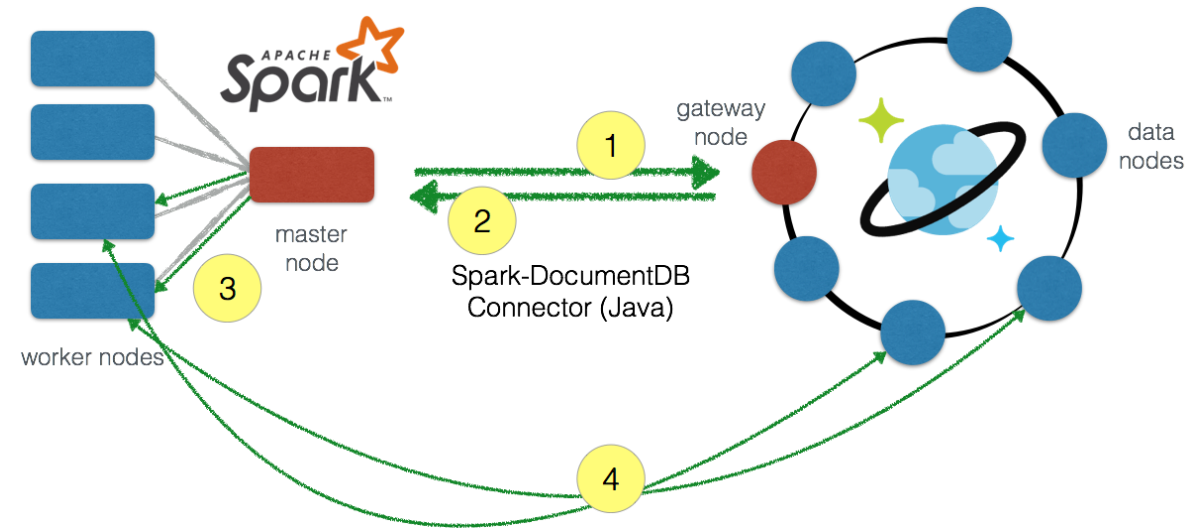
- Embedded into applications
- Strong and deep Java dependencies
- Graph traversal
- Single node

- **SQL Server 2017** and **SQLDB**

- Most of data is relational
- Easy integration with other RDBMS features
 - The Machine Learning example
- Different data workloads integration

- **Azure Cosmos DB**

- Planet scale!
- Multi-master and relaxed consistency models
- Low latency and strong SLA
- Stressing over Graph traversal
- Integration: Apache Spark to Azure Cosmos DB Connector
- Gremlin compat: compatibility with the open-source frameworks recommended by Apache TinkerPop



Q&A

Domande e risposte



Contatti

ROZZANO (MI) – BOLOGNA
ROMA – NAPOLI – GENOVA
TORINO

OverNet Education

Info@OverNetEducation.it

www.OverNetEducation.it

Rozzano (MI) +39 02 365738

Bologna +39 051 269911

www.wpc-overneteducation.it



APPENDIX

Session resources

SQL Server/DB Graph documentation: <https://docs.microsoft.com/en-us/sql/relational-databases/graphs/sql-graph-overview>

SQL Server 2017 Graph blog post: <https://blogs.technet.microsoft.com/dataplatforminsider/2017/04/20/graph-data-processing-with-sql-server-2017/>

Welcome to Azure Cosmos DB: <https://docs.microsoft.com/en-us/azure/cosmos-db/introduction>

DB-Engines Ranking of Graph DBMS: <https://db-engines.com/en/ranking/graph+dbms>

Azure Cosmos DB: Perform graph analytics by using Spark and Apache TinkerPop Gremlin: <https://docs.microsoft.com/en-us/azure/cosmos-db/spark-connector-graph>

Free Book Download O'Reilly's Graph Databases: <https://neo4j.com/graph-databases-book/?ref=home>

SAMPLE: Recommendation System 2 million nodes, 48.5 million edges)

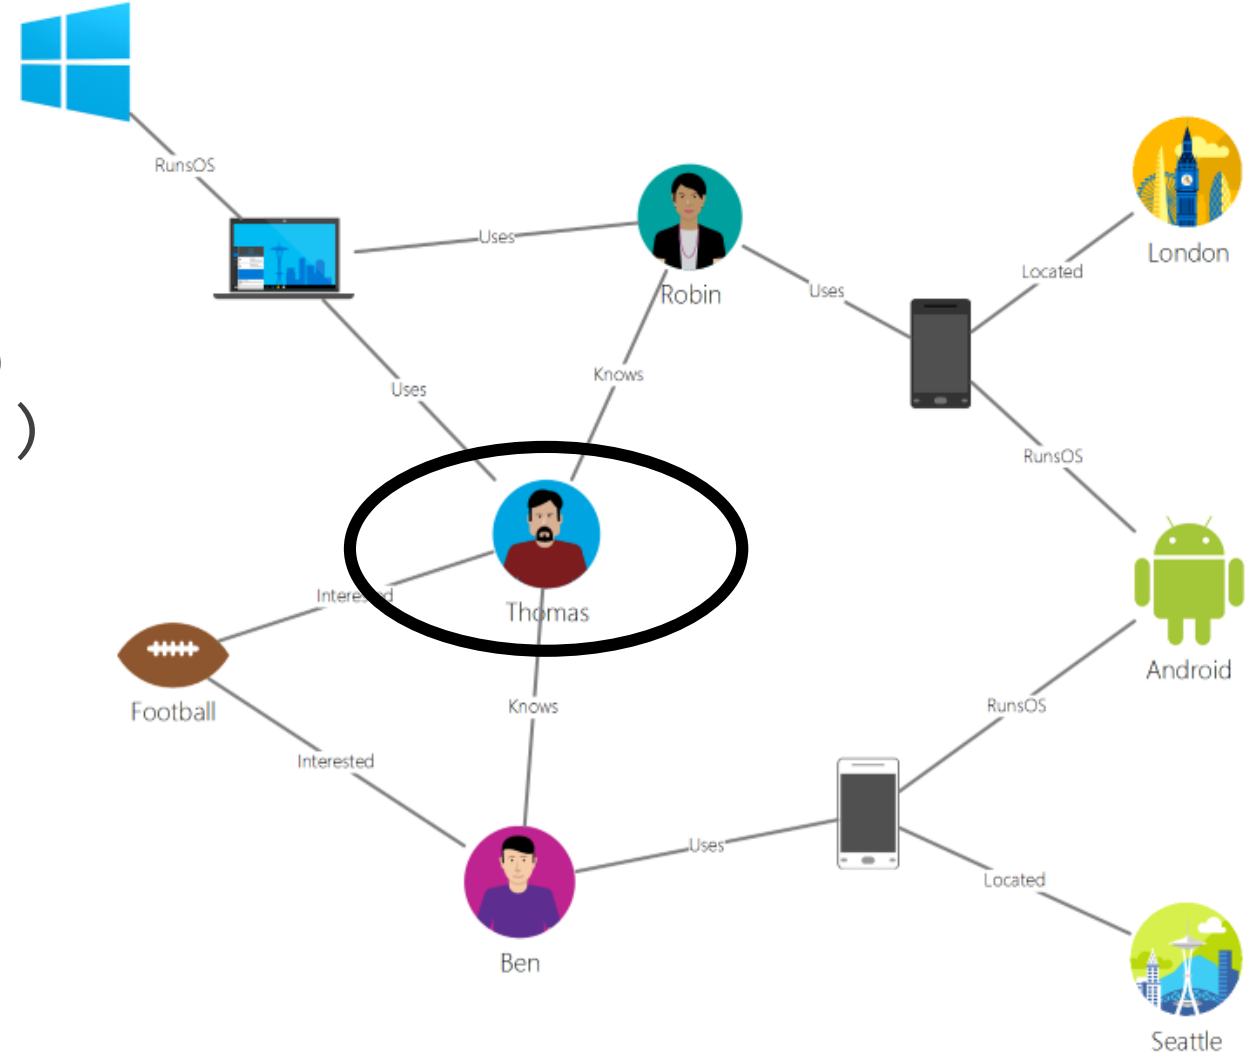
- <https://blogs.msdn.microsoft.com/sqlcat/2017/04/21/build-a-recommendation-system-with-the-support-for-graph-data-in-sql-server-2017-and-azure-sql-db/>
- <https://github.com/arvindshmicrosoft/MillionSongDatasetinSQLServer>

SAMPLE: A more comprehensive end-to-end scenario which features Graph Data as one of the building blocks

- <https://github.com/Microsoft/sql-server-samples/blob/master/samples/applications/iot-connected-car>

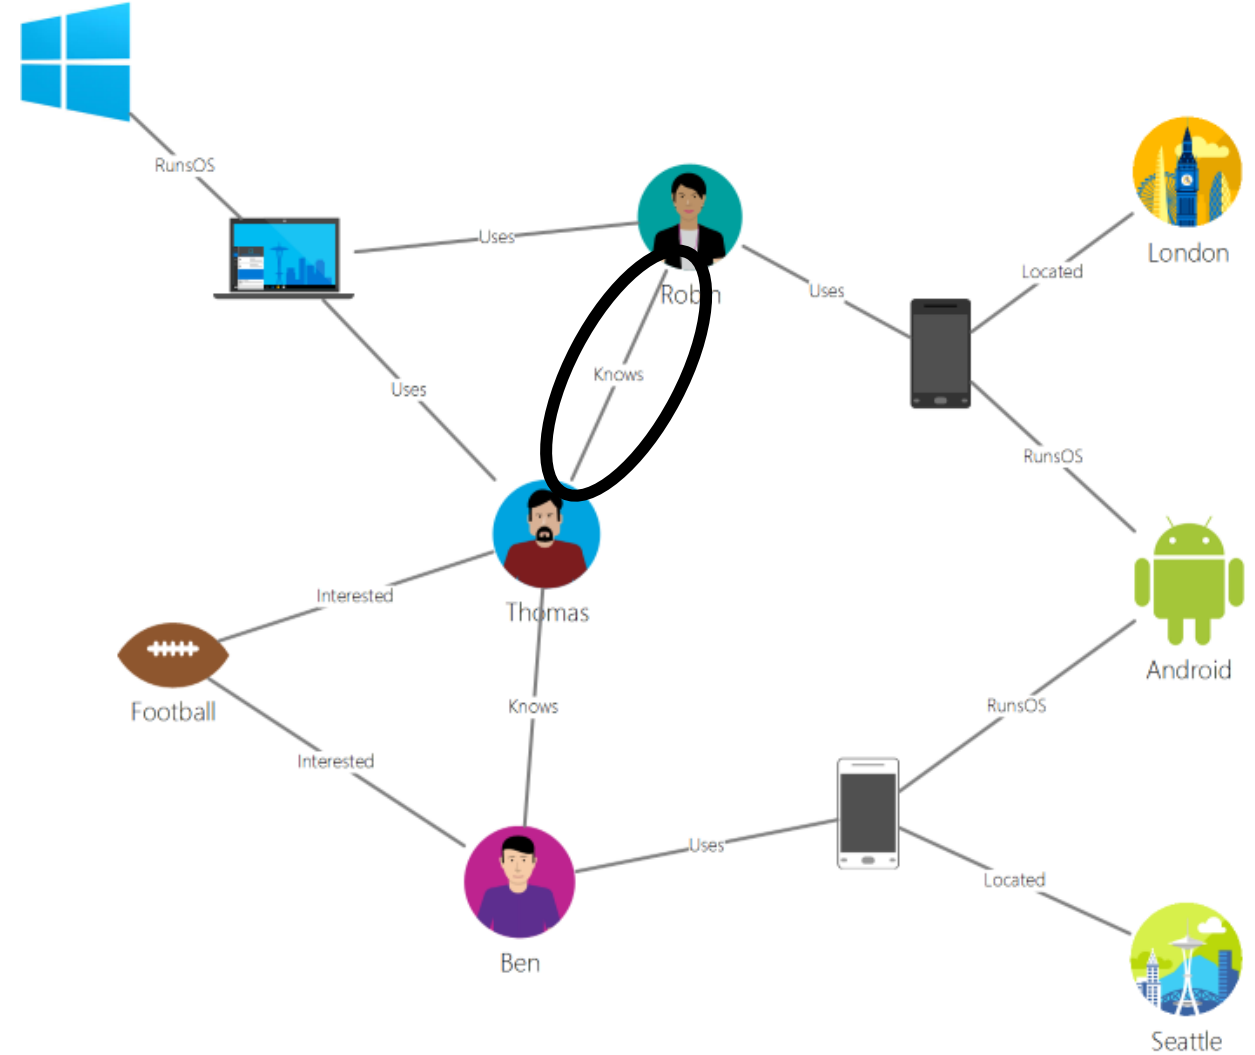
Add vertex

```
g.addV('person')  
  .property('id', 'Thomas')  
  .property('firstName', 'Thomas')  
  .property('lastName', 'Andersen')  
  .property('age', 44)  
  .property('userid', 1)
```



Add edge

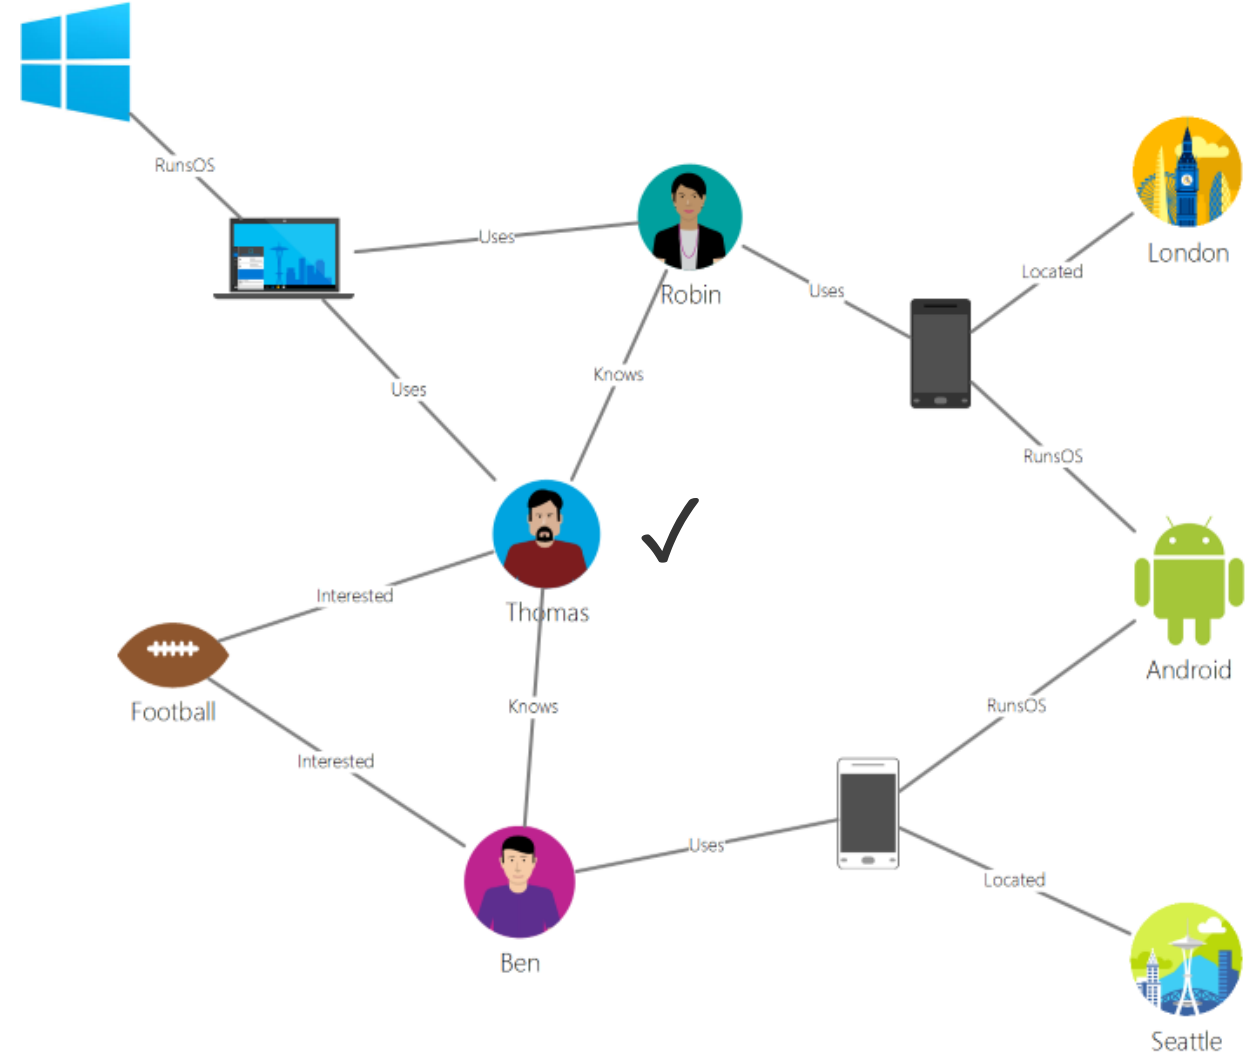
```
g.V('Thomas')
  .addE('knows')
  .property('since', '2017')
  .to(g.V('Robin'))
```



Filter

```
g.V()  
  .hasLabel('person')  
  .has('age', gt(40))
```

- ✓ Automatic index on /label/person
- ✓ Automatic index on /age/[>40]

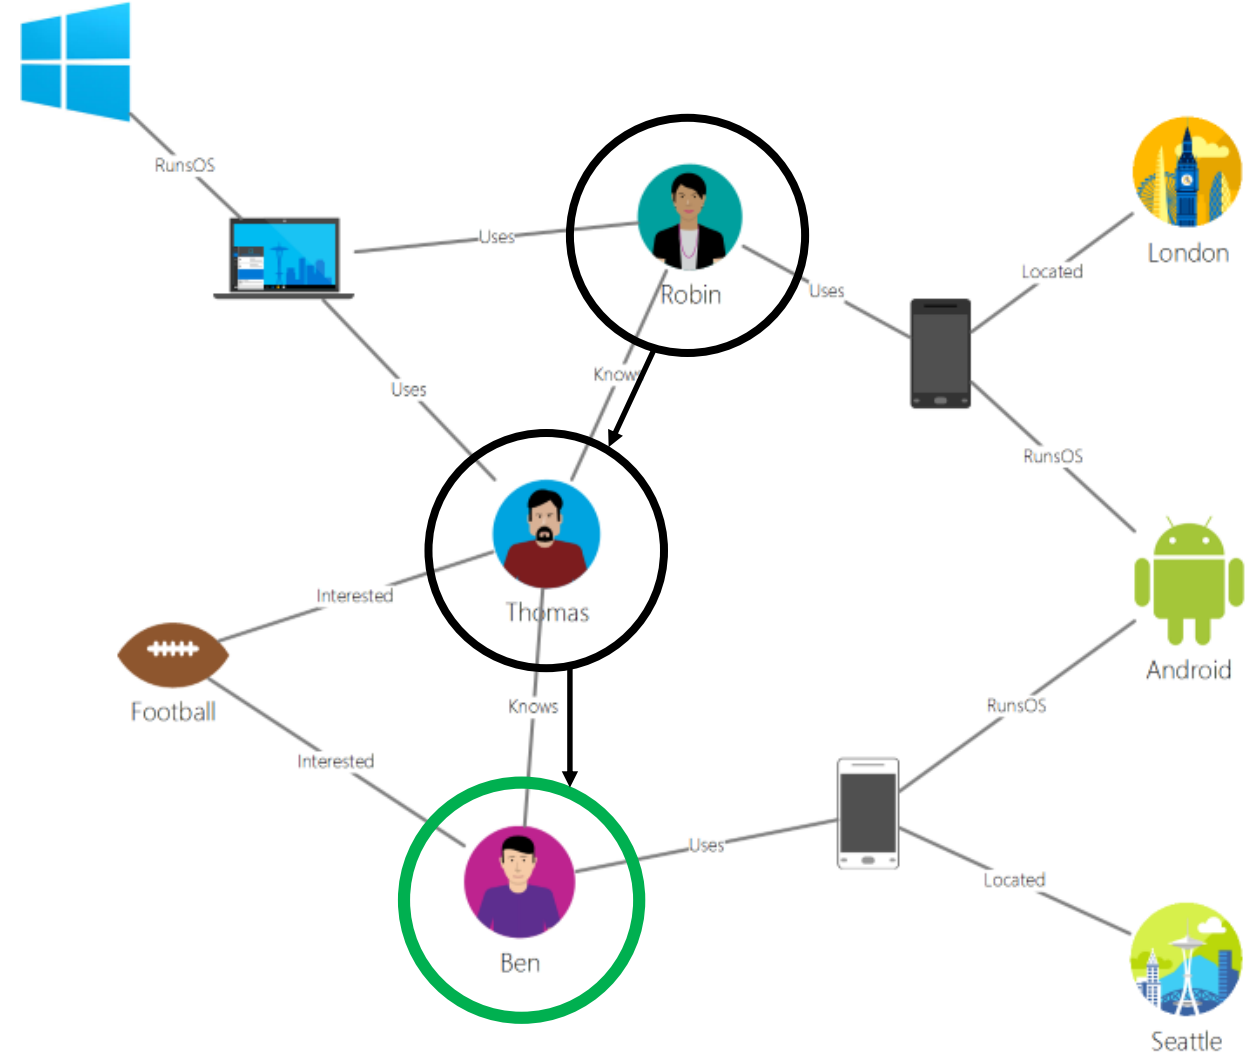


Traverse & Closure

```
g.V('robin')
  .out('knows')
  .hasLabel('person')
```

```
g.V('robin')
  .out('knows')
  .hasLabel('person')
  .out('knows')
  .hasLabel('person')
```

```
g.V('robin')
  .repeat(out())
  .until(has('id', 'ben'))
  .path()
```



From SQL to Gremlin

Projection

```
g.V().hasLabel('person')  
    .values('firstName')
```

Sort

```
g.V().hasLabel('person')  
    .order().by('firstName', decr)
```

Filter

```
g.V().hasLabel('person')  
    .has('age', gt(40))
```

Aggregate

```
g.V().hasLabel('person').count()
```

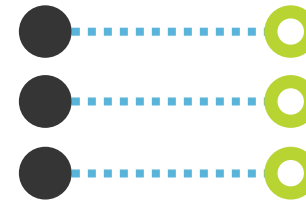
Join

```
g.V('robin').out('knows').hasLabel('person')  
    .out('knows').hasLabel('person')
```

Graph data format (multi-model)

Who wants to have 4 different databases?

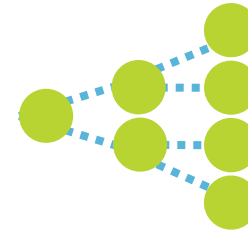
- Database engine operates on atom-record-sequence (ARS) based type system
- All data models are efficiently translated to ARS
- API and wire protocols are supported via extensible modules
- Instance of a given data model can be materialized as trees



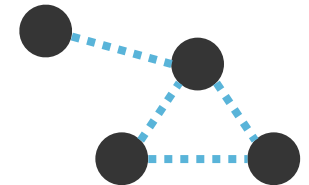
KEY-VALUE



COLUMN-FAMILY



DOCUMENT



GRAPH

Graph compute tier

- Two-tier architecture
- Independently scalable graph compute tier for graph
- Data tier with low latency, automatic indexing, global distribution, etc.

